

Distributed Systems Design

COMP 6231

Architectures

Lecture 2

Essam Mansour

Today...

- Last Session:

- Introduction

- This Session:

- Architectures for distributed systems (*Chapter 2*)
 - Architectural styles
 - Client-server architectures
 - Decentralized and peer-to-peer architectures



Signal ✓
@signalapp

Signal is experiencing technical difficulties. We are working hard to restore service as quickly as possible.

11:33 AM · 2021-01-15 · [Twitter Web App](#)

3,311 Retweets 1,387 Quote Tweets 22.6K Likes



Signal ✓ @signalapp · 2h
Replying to @signalapp

We have been adding new servers and extra capacity at a record pace every single day this week nonstop, but today exceeded even our most optimistic projections. Millions upon millions of new users are sending a message that privacy matters. We appreciate your patience.

405

3,359

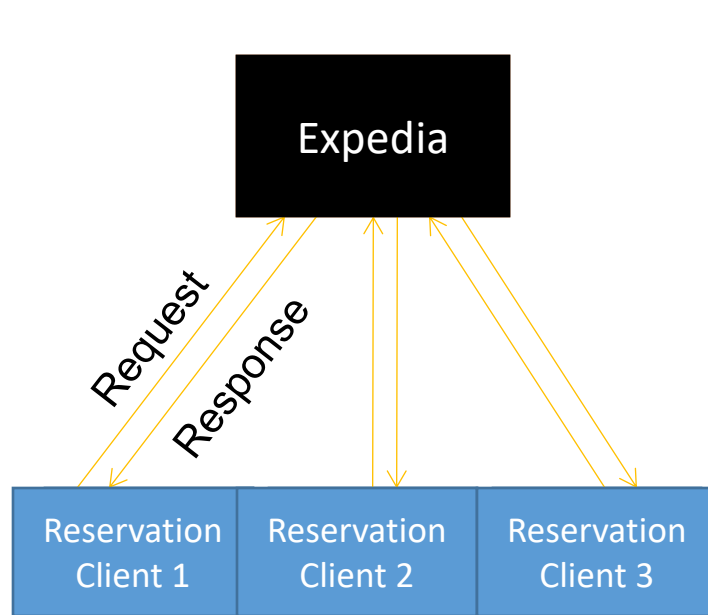
13.7K




Signal ✓ @signalapp · 1h

We are making progress towards getting the service back online. Privacy is our top priority, but adding capacity is a close

Bird's Eye View of Some Distributed Systems



Google Search
Airline Booking



Bit-torrent
Skype

Is it still
the case
?

How would one characterize these distributed systems?


Simple Characterization of Distributed Systems

- What are the entities that are communicating in a DS?
 - a) Communicating entities (system-oriented vs. problem-oriented entities)
- How do the entities communicate?
 - b) Communication paradigms (sockets and RPC)
- What roles and responsibilities do the entities have?
 - c) This could lead to different organizations (referred, henceforth, to as *architectures*)

Architectural Styles

- Important styles of architecture for distributed systems
 - Layered architectures
 - Object-based architectures
 - Data-centered architectures
 - Event-based architectures
 - Resource-based architectures


Layered Design



- Each layer uses previous layer to implement new functionality that is exported to the layer above
- Example: Multi-tier web apps

Layering

- A complex system is partitioned into layers
 - Upper layer utilizes the services of the lower layer
 - *A vertical organization of services*
- For example, a three-layer solution could easily be deployed on a single tier, such as a personal workstation.
- Layering simplifies the design of complex distributed systems by hiding the complexity of below layers
- Control flows from layer to layer



Layering – Platform and middleware

- Distributed systems can be organized into three layers:

1. Platform

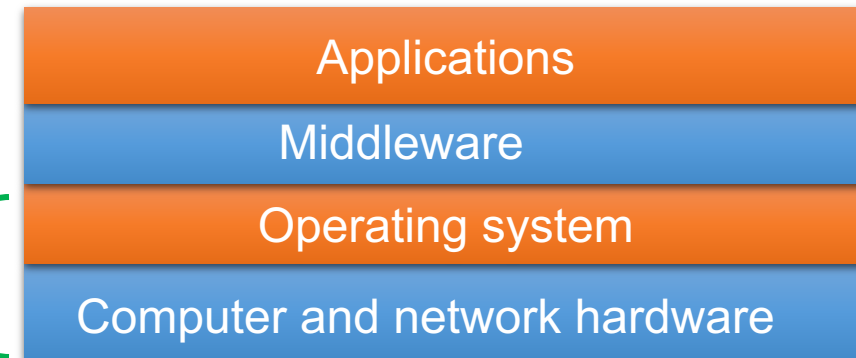
- Low-level hardware and software layers
- Provides common services for higher layers

2. Middleware


- Masks heterogeneity and provides convenient programming models to application programmers
- Typically, it simplifies application programming by abstracting communication mechanisms

3. Applications

Platform



Object-based Style

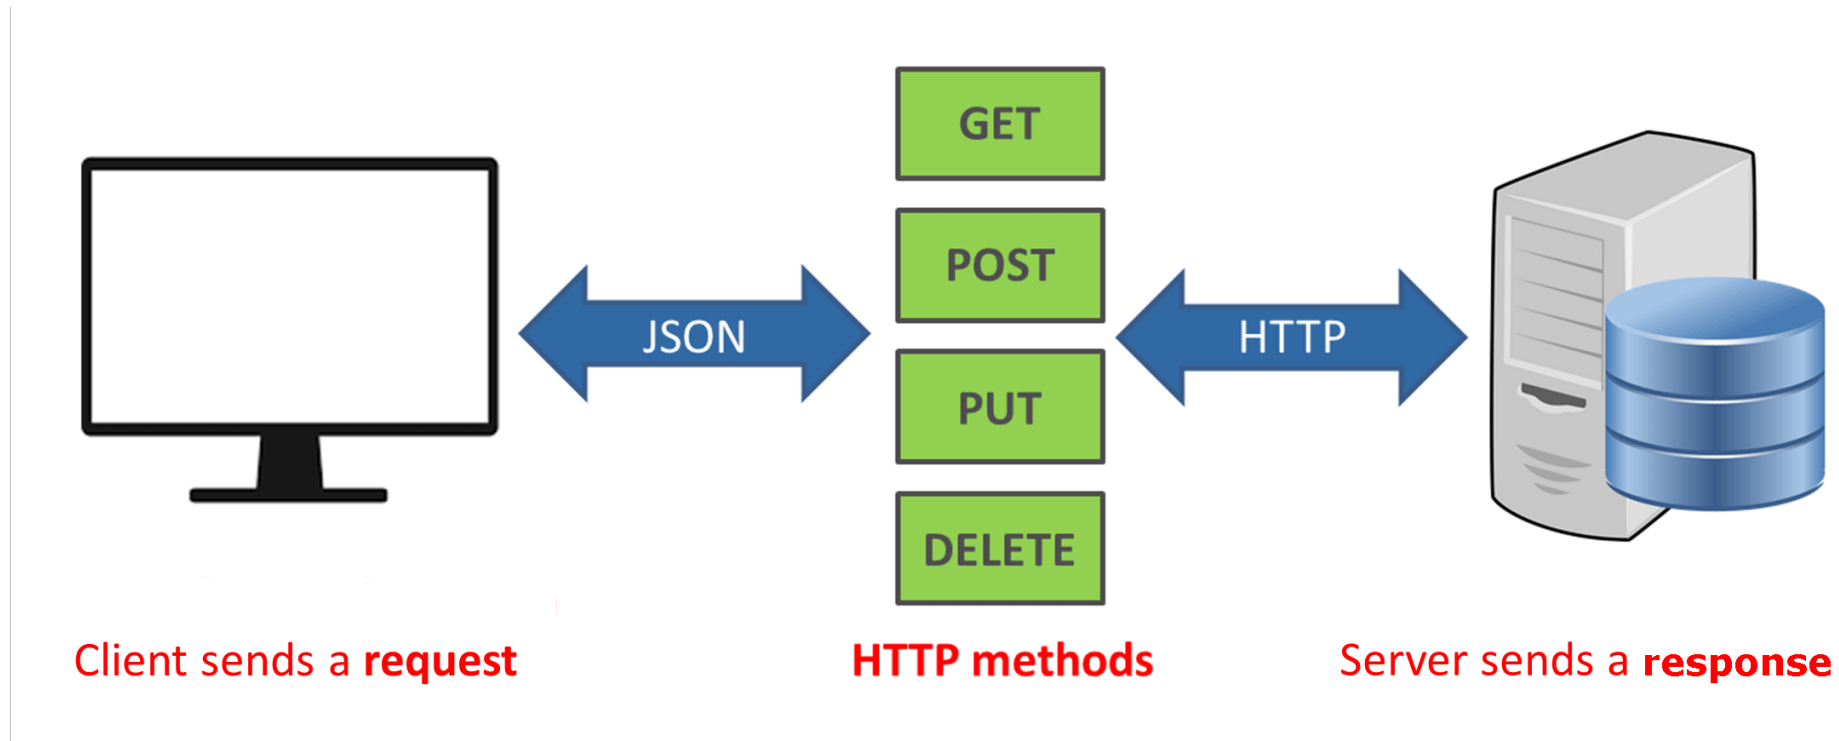


(b)

- Each object corresponds to a components
- Components interact via remote procedure calls
 - Popular in client-server systems

Resource-oriented Architecture

- Example of ROA: Representational State Transfer (REST)




Resource-oriented Architecture

- Example of ROA: Representational State Transfer (REST)

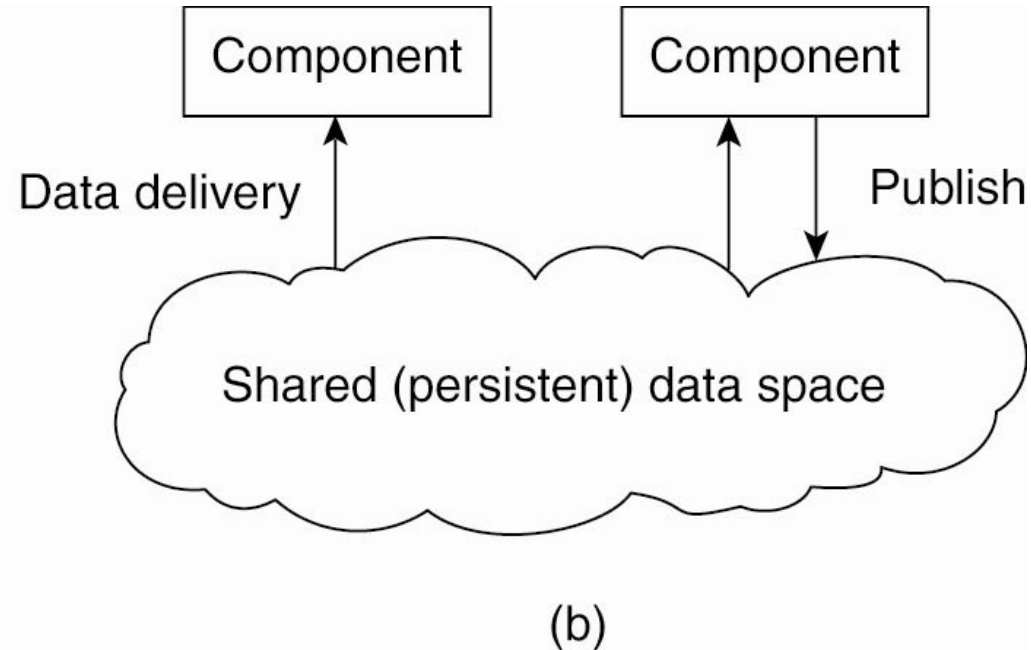
- Basis for RESTful web services
- Resources identified through a single naming scheme
 - Uniform Resource Identifier (URI)
- All services offer same interface (e.g., 4 HTTP operations)
 - **Get** / **Put** / **Delete** / **Post** HTTP operations
- Messages are fully described
- No state of the caller is kept (stateless execution)
- Example: use HTTP for API
 - <http://bucketname.s3.aws.com/objName>
- Return JSON objects
 - ```
{"name": "test.com", "messages": ["msg 1", "msg 2", "msg 3"], "age": 100}
```
- **Discuss:** Service-oriented (SOA) vs. Resource-oriented (ROA)

# Event-based architecture




- Communicate via a common repository
  - Use a publish-subscribe paradigm
  - Consumers subscribe to types of events
  - Events are delivered once published by any publisher

# Shared data-space



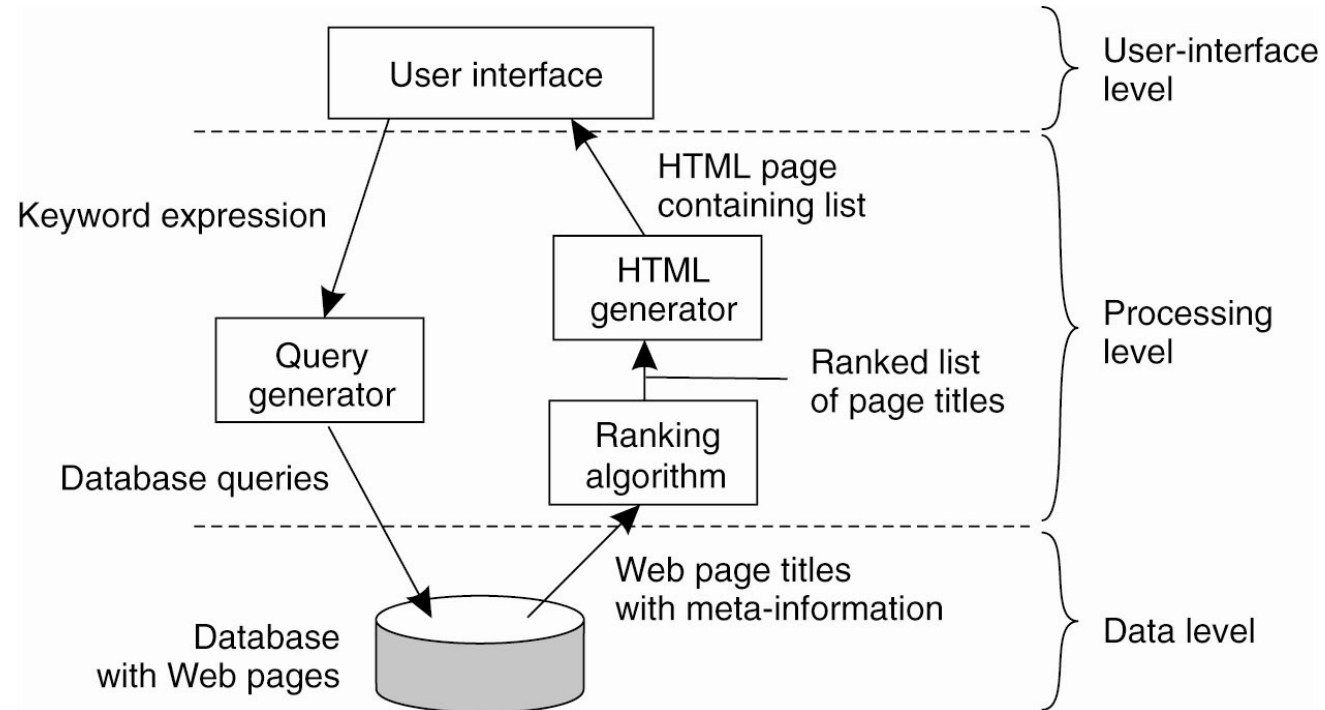
- “Bulletin-board” architecture
  - Decoupled in space and time
  - Post items to shared space; consumers pick up at a later time

# Client-Server Architectures



- Most common style: client-server architecture
- Application layering
  - User-interface level
  - Processing level
  - Data level

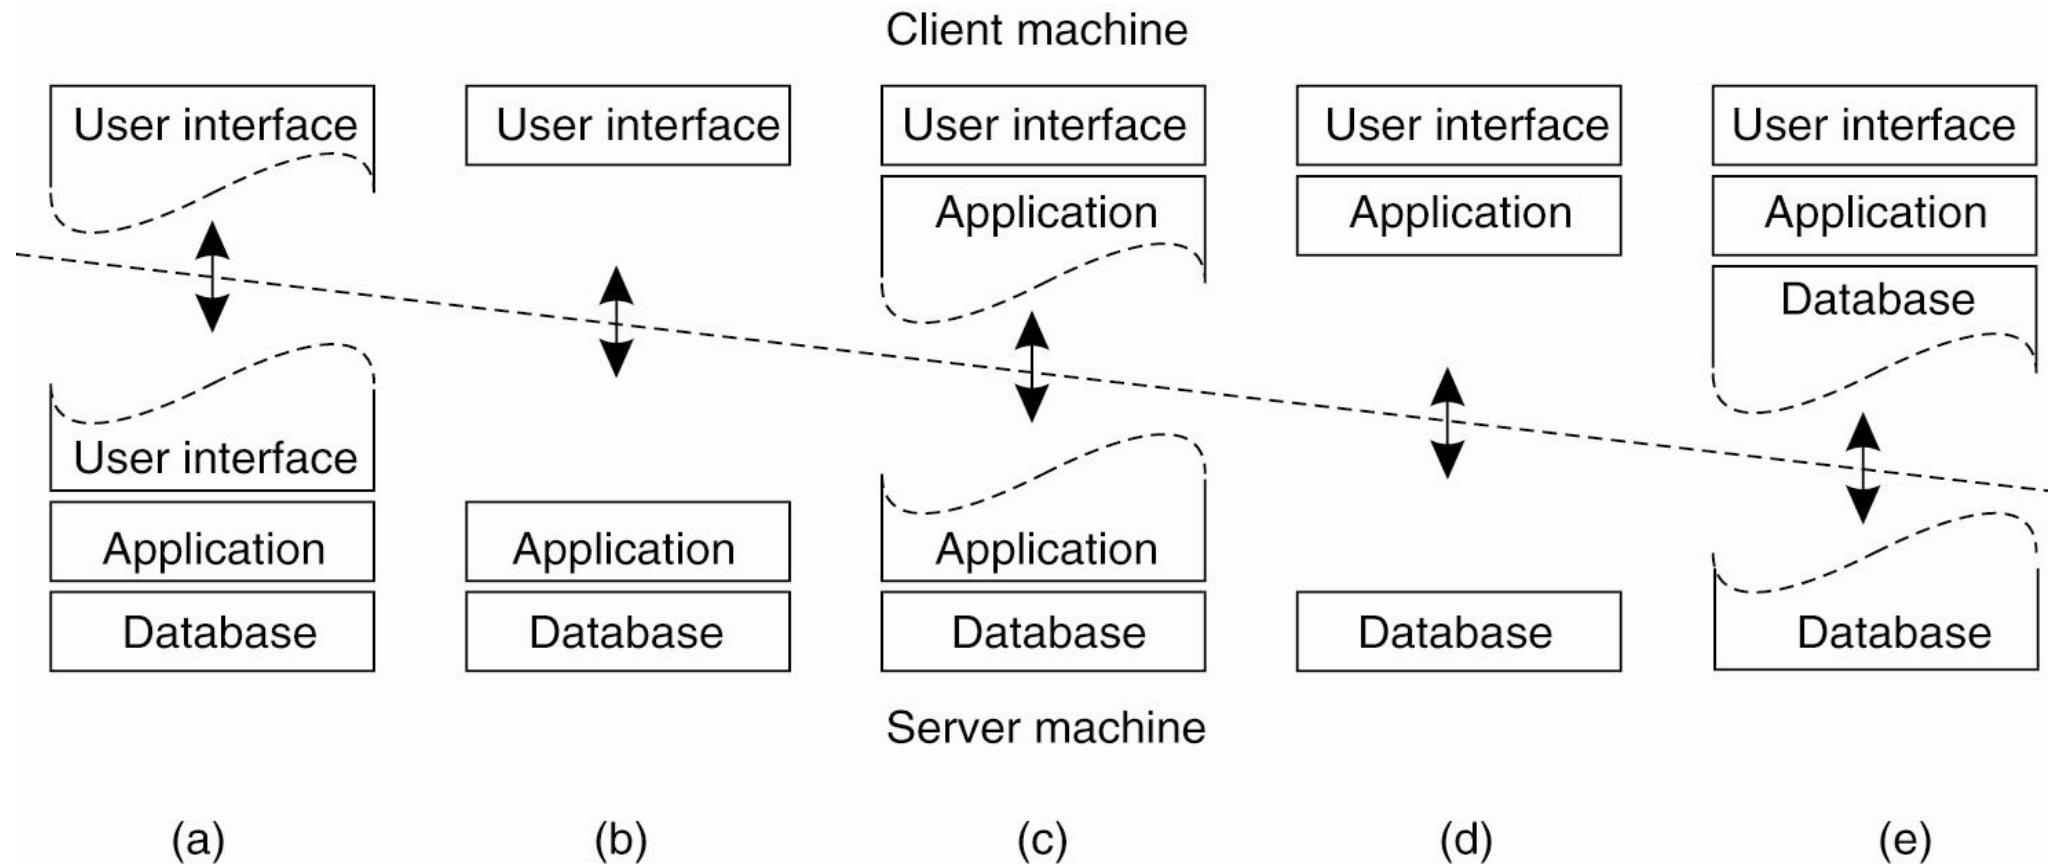
# Search Engine Example



Search engine architecture with 3 layers




# A Spectrum of Choices



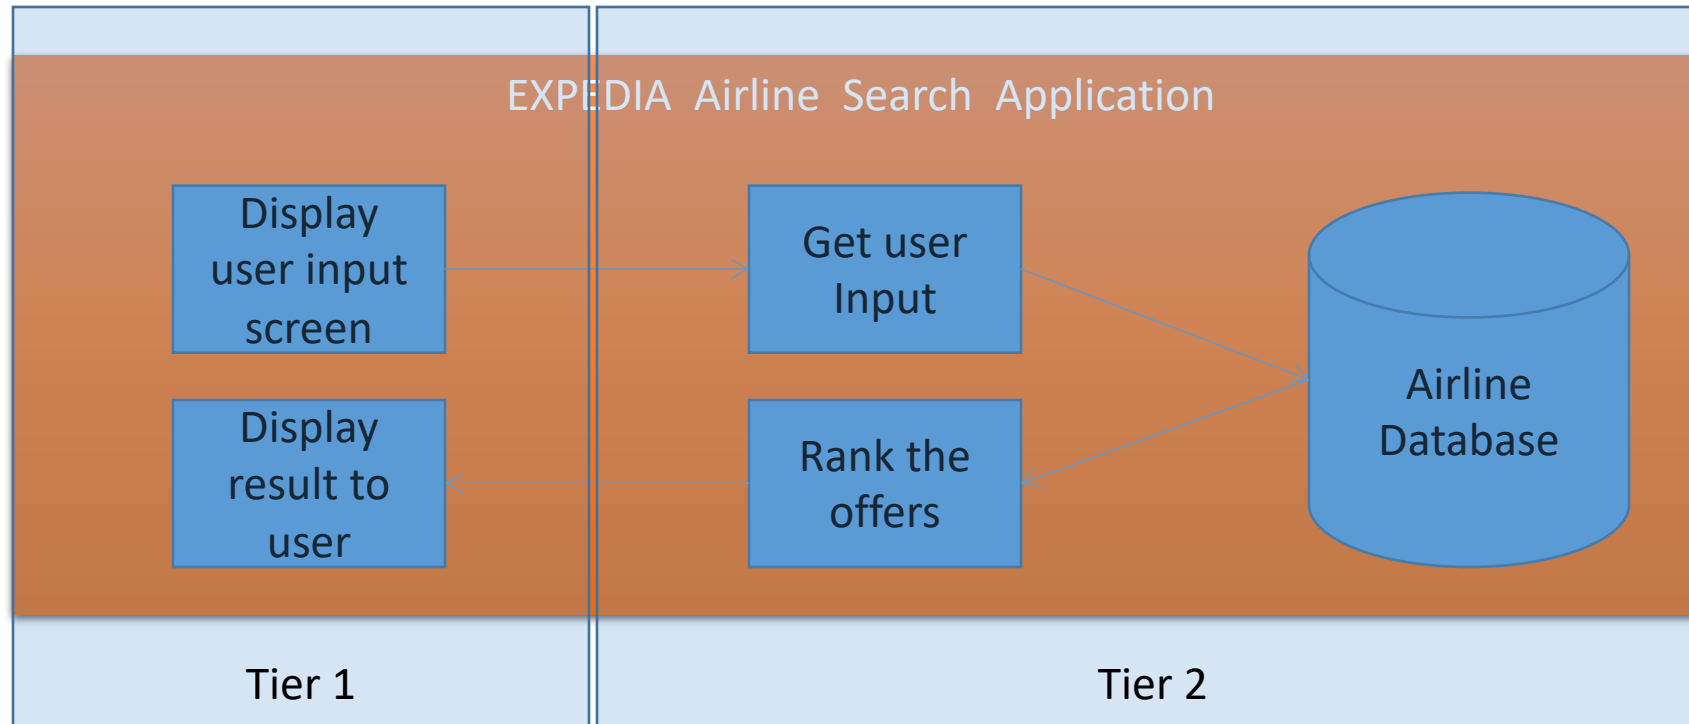
# Tiering

- Tiering is a technique to:
  1. Organize the functionality of a service,
  2. and place the functionality into appropriate servers
  3. a tier is a physical structuring mechanism for the system infrastructure



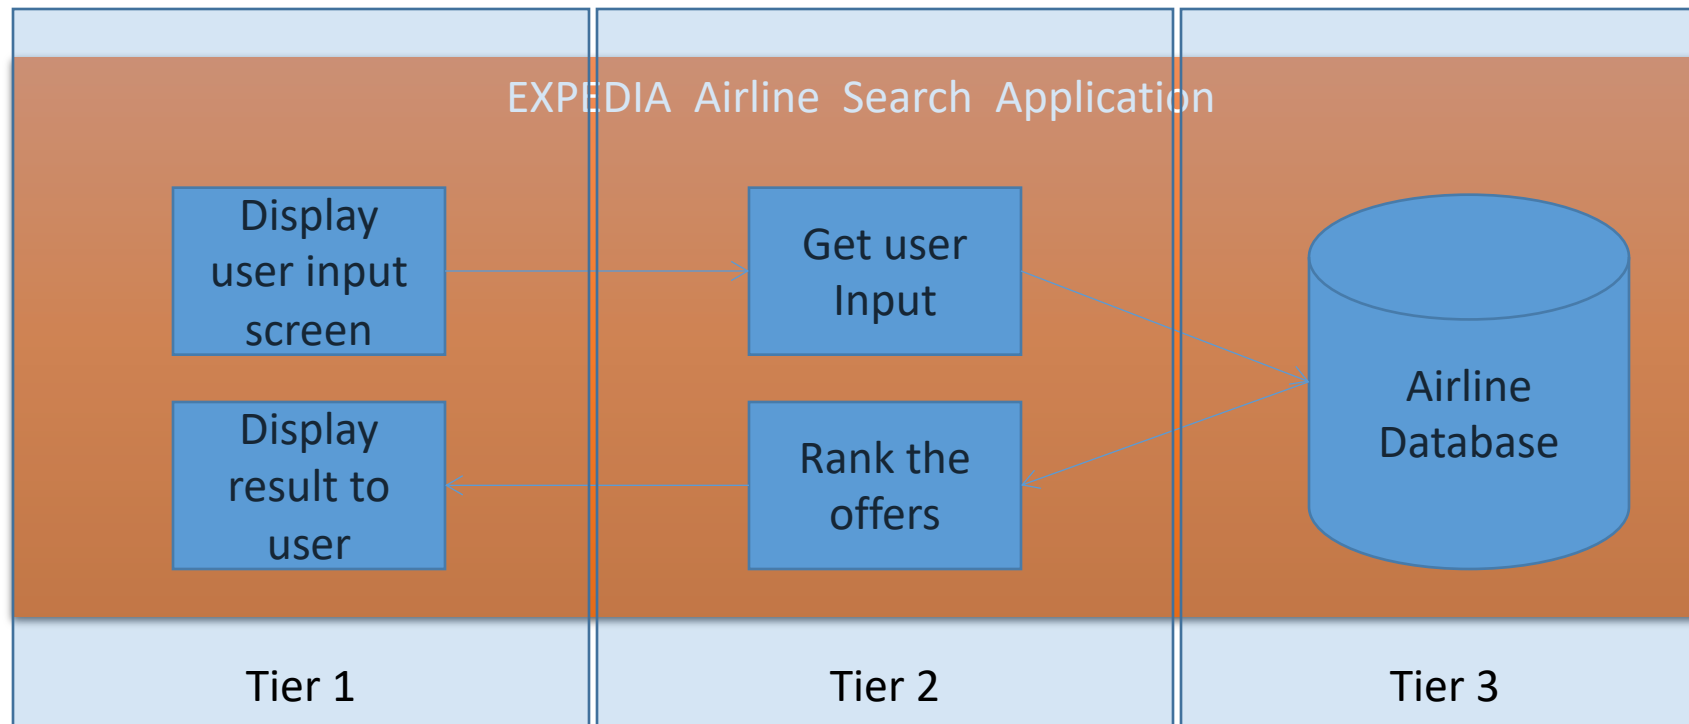
# A Two-Tiered Architecture

- How would you design an airline search application?

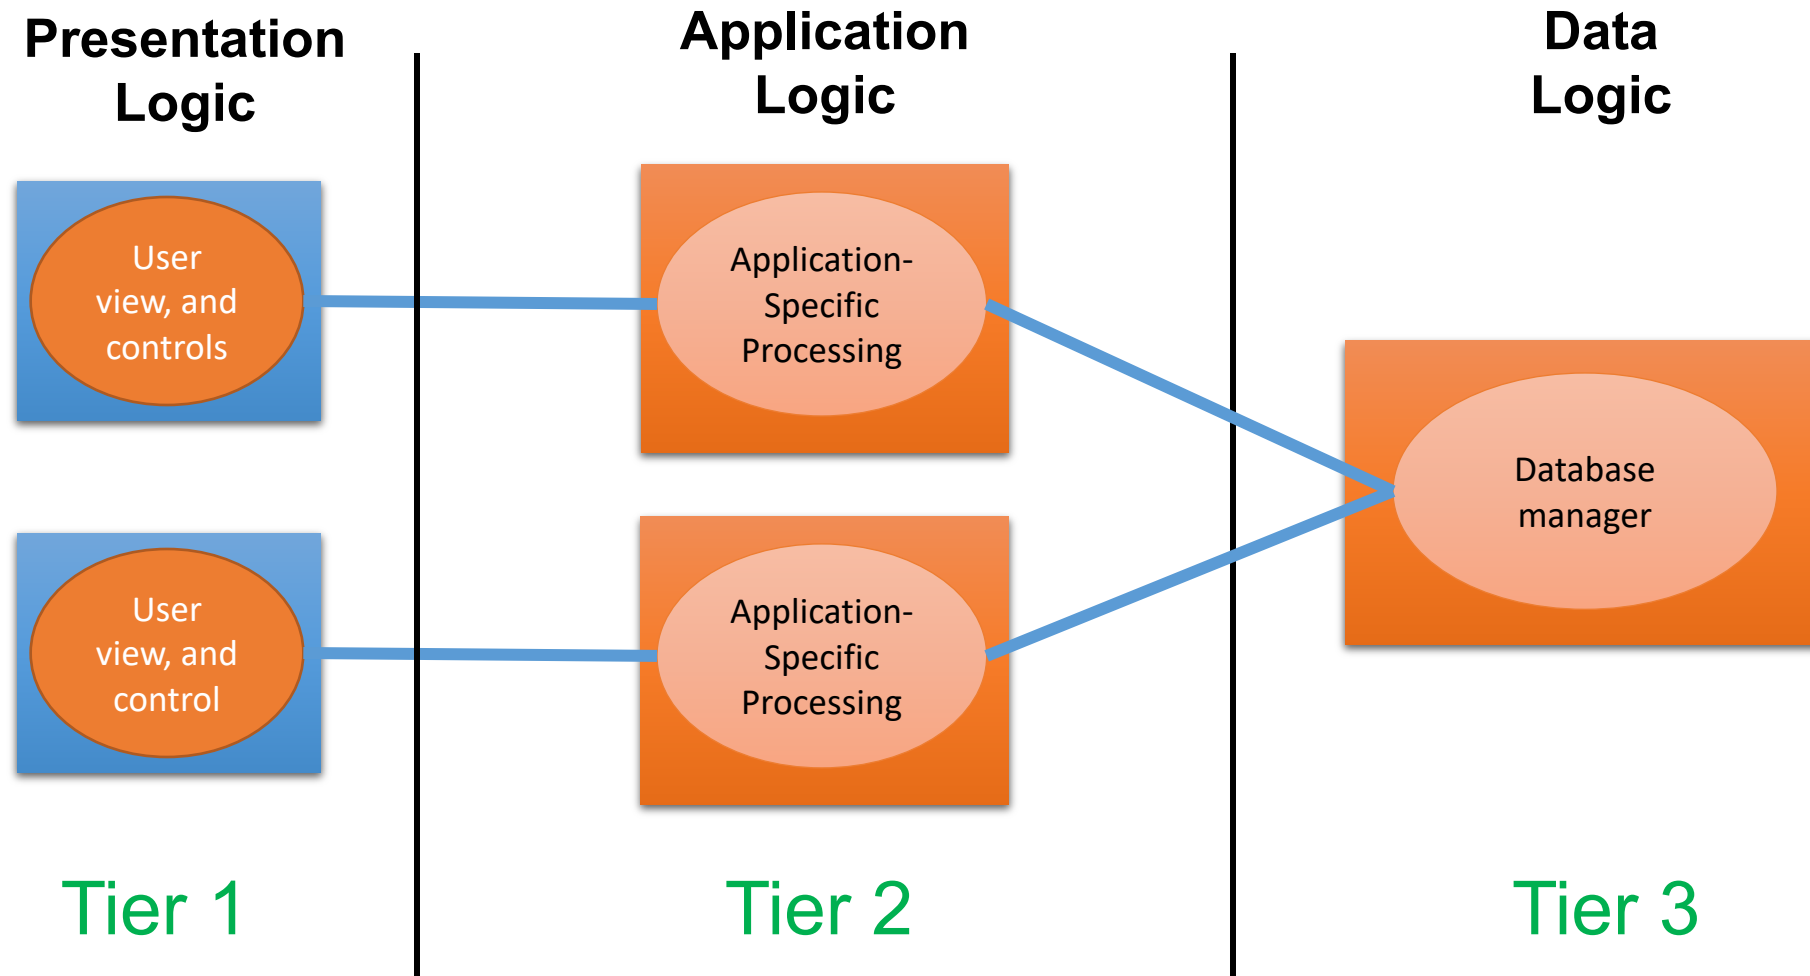


# A Three-Tiered Architecture

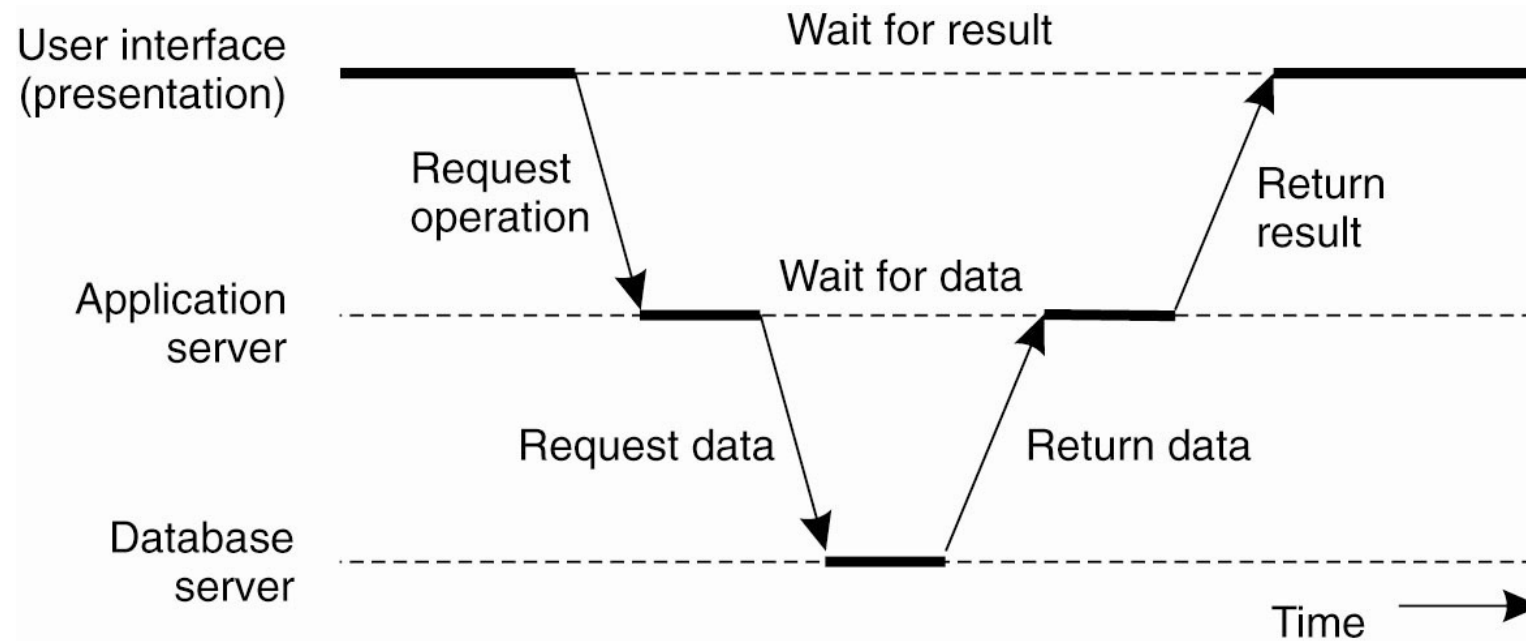
- How would you design an airline search application?



# A Three-Tiered Architecture



# Three-tier Web Applications



- Server itself uses a “client-server” architecture
- 3 tiers: HTTP, J2EE and database
  - Very common in most web-based applications

# Three-Tiered Architecture: Pros and Cons


- Advantages:

- Enhanced maintainability of the software (one-to-one mapping from logical elements to physical servers)
- Each tier has a well-defined role

- Disadvantages:

- Added complexity due to managing multiple servers
- Added network traffic
- Added latency


# Edge-Server Systems



- Edge servers: from *client-server* to *client-proxy-server*
- Content distribution networks: proxies cache web content near the edge




# Edge-Server Systems




[https://en.wikipedia.org/wiki/Edge\\_computing](https://en.wikipedia.org/wiki/Edge_computing)

# Decentralized Peer-to-Peer Architecture

- A peer-to-peer (P2P) architecture can be characterized as follows:
  - 1) All nodes are equal (no hierarchy)
    - No Single-Point-of-Failure (SPOF)
  - 2) A central coordinator is not needed
    - But, decision making becomes harder
  - 3) The underlying system can scale out indefinitely
    - In principle, no performance bottleneck



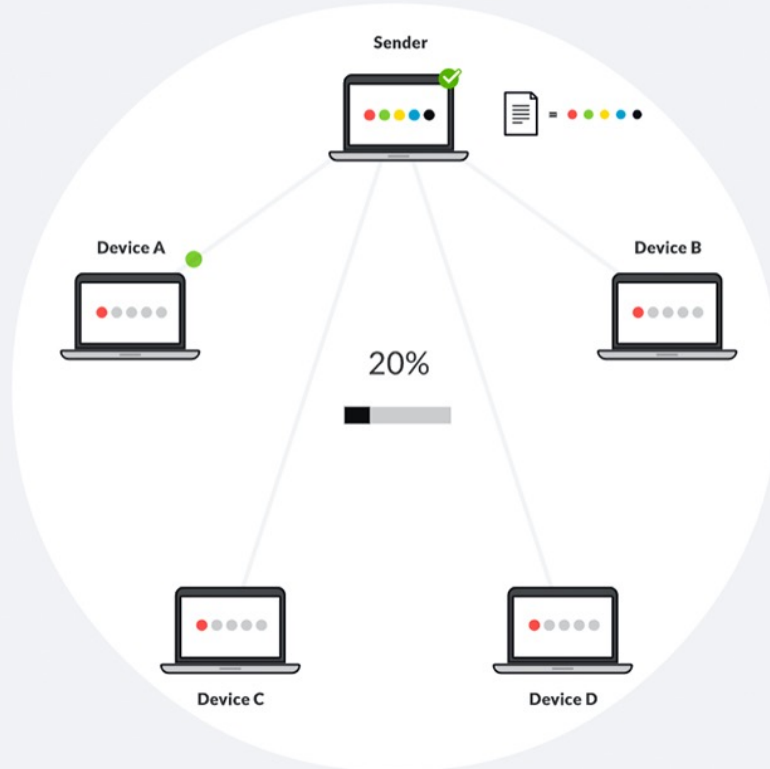
# Decentralized Peer-to-Peer Architecture



- Peer-to-peer systems
  - Removes distinction between a client and a server
  - Overlay network of nodes
- Chord: structured peer-to-peer system
  - Use a distributed hash table to locate objects
    - Data item with key  $k$  -> smallest node with  $id \geq k$

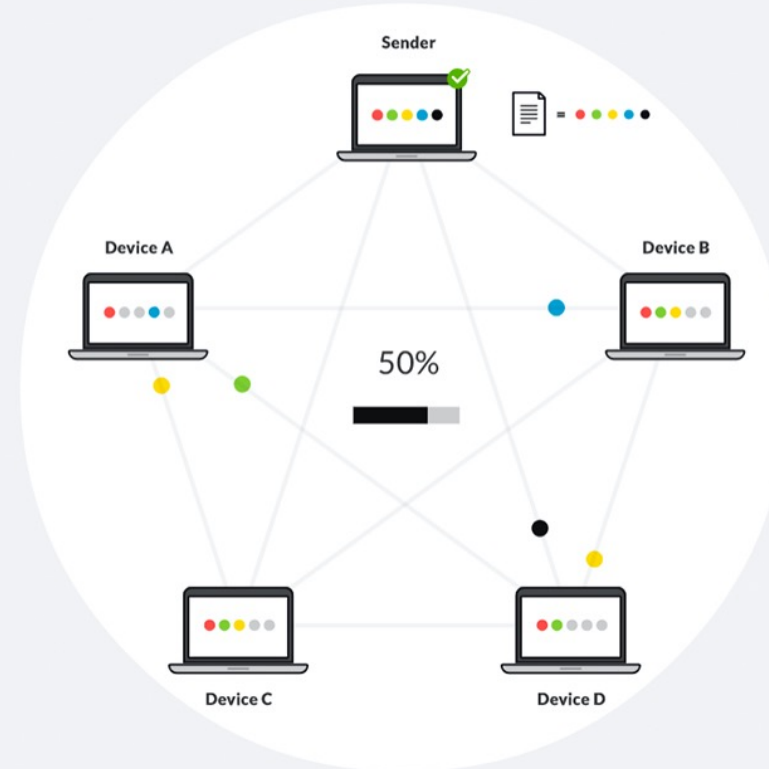
# Client-Server vs Peer-to-Peer

**Client-Server**  
Architecture



Vs.

**Peer-to-Peer**  
Architecture

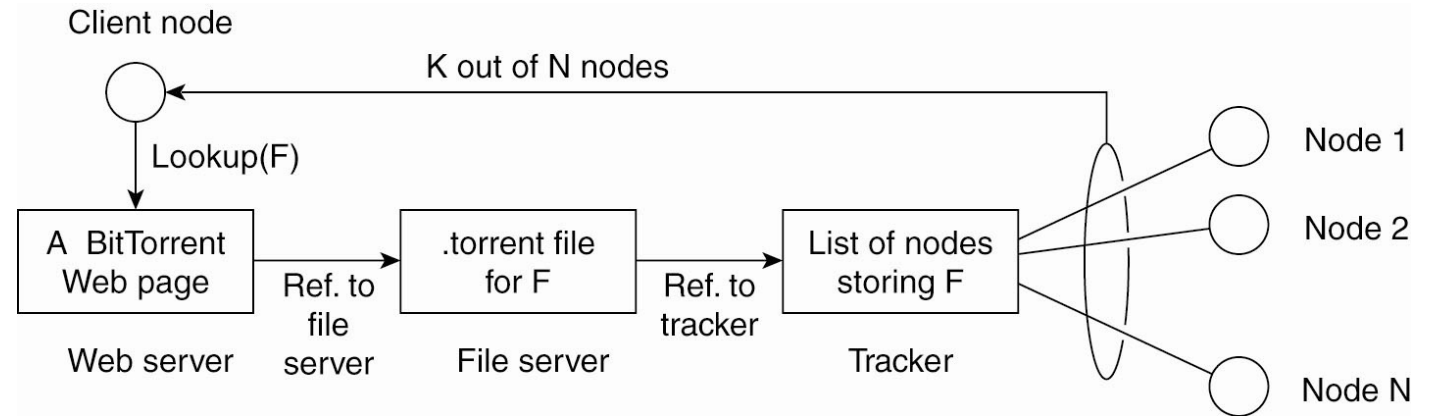


# Peer-to-Peer Architecture

- A peer-to-peer (P2P) architecture can be characterized as follows:
  - 4) Peers can interact directly, forming groups and sharing contents (or offering services to each other)
    - At least one peer should share the data, and this peer should be accessible
    - Popular data will be highly available (it will be shared by many)
    - Unpopular data might eventually disappear and become unavailable (as more users/peers stop sharing them)
  - 5) Peers can form a virtual *overlay network* on top of a physical network topology
    - *Logical paths* do not usually match *physical paths* (i.e., higher latency)
    - Each peer plays a role in routing traffic through the overlay network


# Peer-to-Peer Architecture

## Collaborative Distributed Systems




- **BitTorrent: Collaborative P2P downloads**
  - Download chunks of a file from multiple peers
    - Reassemble file after downloading
  - Use a global directory (web-site) and download a .torrent
    - .torrent contains info about the file
      - Tracker: server that maintains active nodes that have requested chunks
      - Force altruism:
        - If  $P$  sees  $Q$  downloads more than uploads, reduce rate of sending to  $Q$


# P2P Types




# P2P Types



Hybrid (Napster)



Unstructured overlay



Structured overlay




# P2P Types

## ■ Unstructured P2P:

- The architecture does not impose any particular structure on the overlay network
  - Each node pick a random set of nodes and becomes their neighbors
  - Choice of degree impacts network dynamics
- Advantages:
  - Easy to build
  - Highly robust against high rates of churn (i.e., when a great deal of peers frequently join and leave the network)
- Main disadvantage:
  - Peers and contents are *loosely-coupled*, creating a data location problem
    - Searching for data might require broadcasting

# P2P Types



# P2P Types

- **Structured P2P:**

- The architecture imposes some structure on the overlay network topology


- **Main advantage:**

- Peers and contents are *tightly-coupled* (e.g., through hashing), simplifying data location

- **Disadvantages:**

- Harder to build
- For optimized data location, peers must maintain extra metadata (e.g., lists of neighbors that satisfy specific criteria)
- Less robust against high rates of churn

# P2P Types



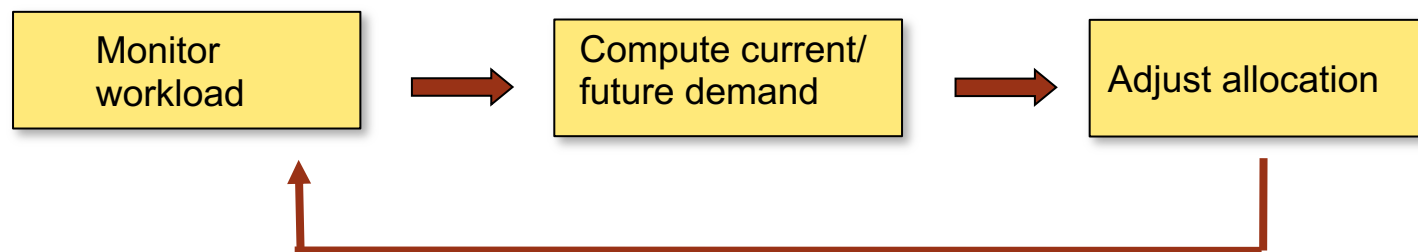
# P2P Types

- Hybrid P2P:

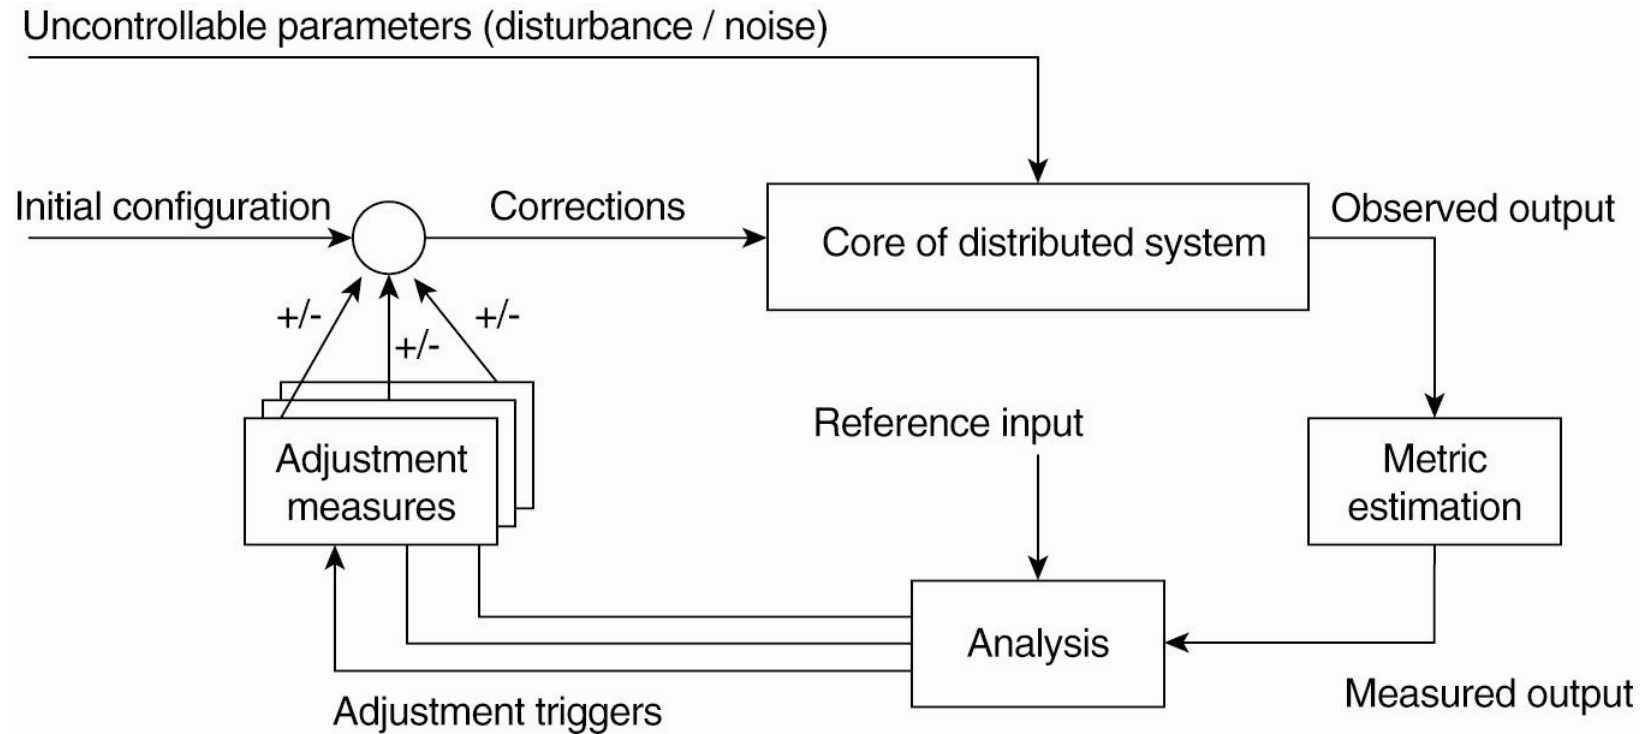
- The architecture can use *some* central servers to help peers locate each other
  - A combination of P2P and master-slave models
- It offers a trade-off between the *centralized functionality* provided by the master-slave model and the *node equality* afforded by the *pure* P2P model
  - In other words, it combines the advantages of the master-slave and P2P models and precludes their disadvantages

# Self-Managing Systems

- System is adaptive
  - Monitors itself and takes action autonomously when needed
    - Autonomic computing, self-managing systems
- Self-\*: self-managing, self-healing
- Example: automatic capacity provisioning
  - Vary capacity of a web server based on demand



# Feedback Control Model



- Use feedback and control theory to design a self-managing system

# A To-Do List

- Read Chapters 2, Architectures
- Project teams