

Distributed Systems Design

COMP 6231

Processes

Lecture 3 – Part 1

Essam Mansour

Today...


- Last Session:
 - Architectures for distributed systems
- Today's Session:
 - Processes
 - Threads
 - Virtualization

Communicating Entities in Distributed Systems

- Communicating entities in distributed systems can be classified into two types:
 - System-oriented entities
 - Processes
 - Threads
 - Nodes
 - Problem-oriented entities
 - Objects (in *object-oriented programming* based approaches)

Classification of Communication Paradigms

- Communication paradigms can be categorized into *three types* based on *where the entities reside*. If entities are running on:



Processes vs Threads

- A program to be executed needs more than just binary code.
- It needs:
 - Memory, and
 - operating system resources
 - a register may hold an instruction, or storage address,
 - program counter keeps track of where a computer is in its program sequence,
 - “stack” is a data structure that stores information about the active subroutines of a computer program

Processes vs Threads

- A computer process



May we have
Multiple processes
Of the same program?



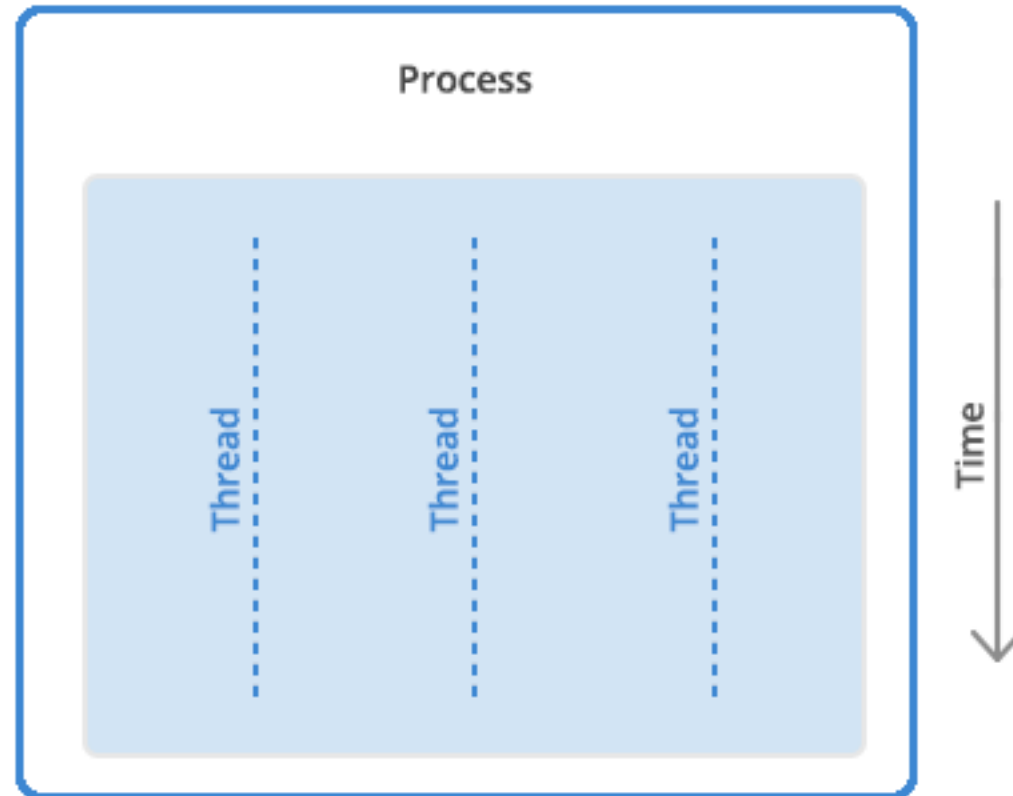
Processes vs Threads

- A computer process
- Yes, but each process has a separate memory address space, i.e., a process runs independently and is isolated from other processes.



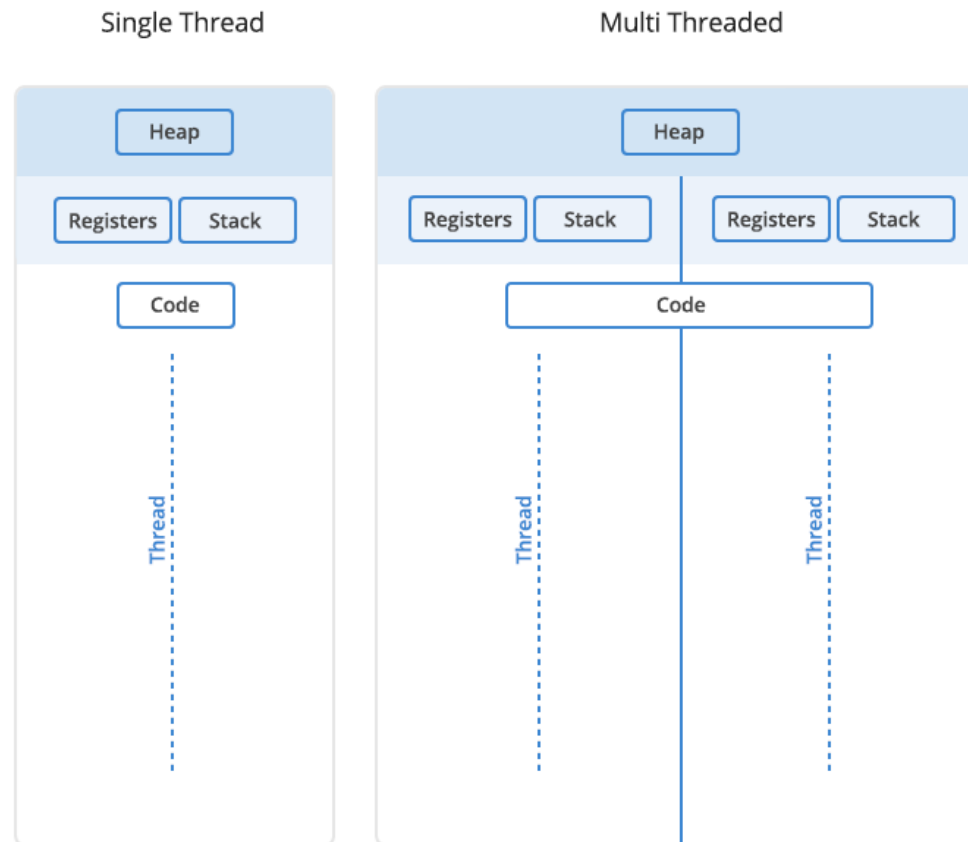
Processes vs Threads

- A thread could be defined as the unit of execution within a process.



Processes vs Threads

- A thread could be defined as the unit of execution within a process.

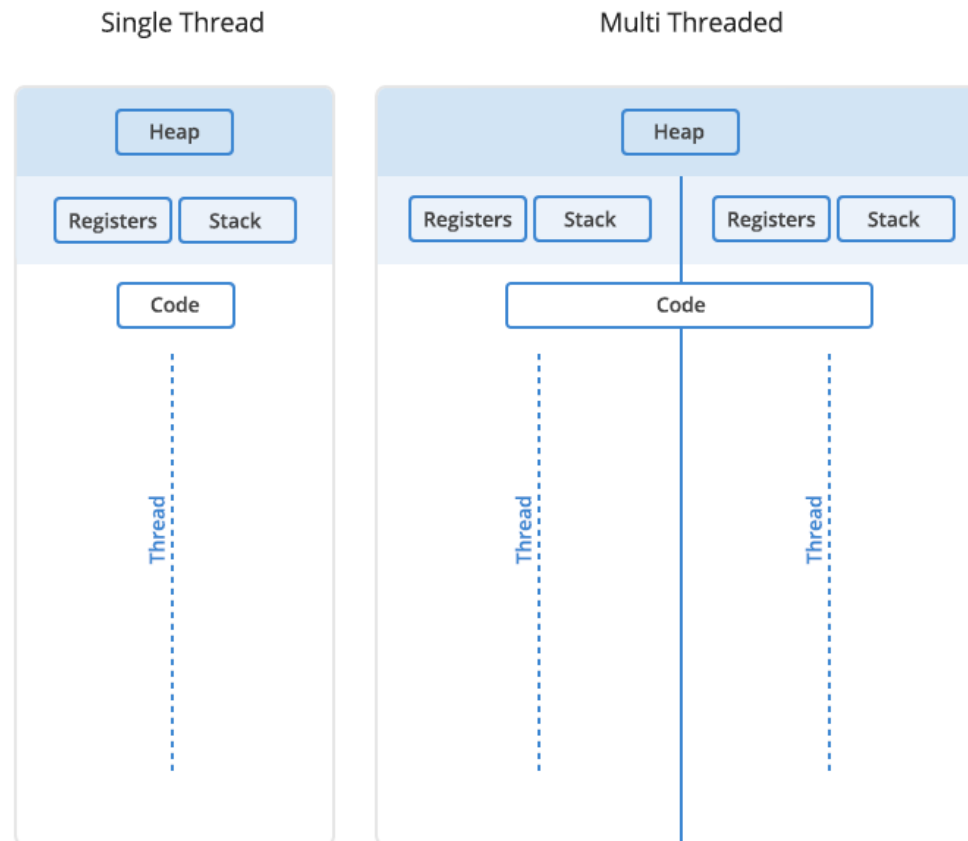


What multithreaded process does share?



Processes vs Threads

- A thread could be defined as the unit of execution within a process.



Threads could be seen as lightweight processes as they have their own stack but can access shared data.

Is it easier to enable communication between threads or Processes?



Processes vs Threads




Figure 3-1. **Context switching** as the result of interprocess communication

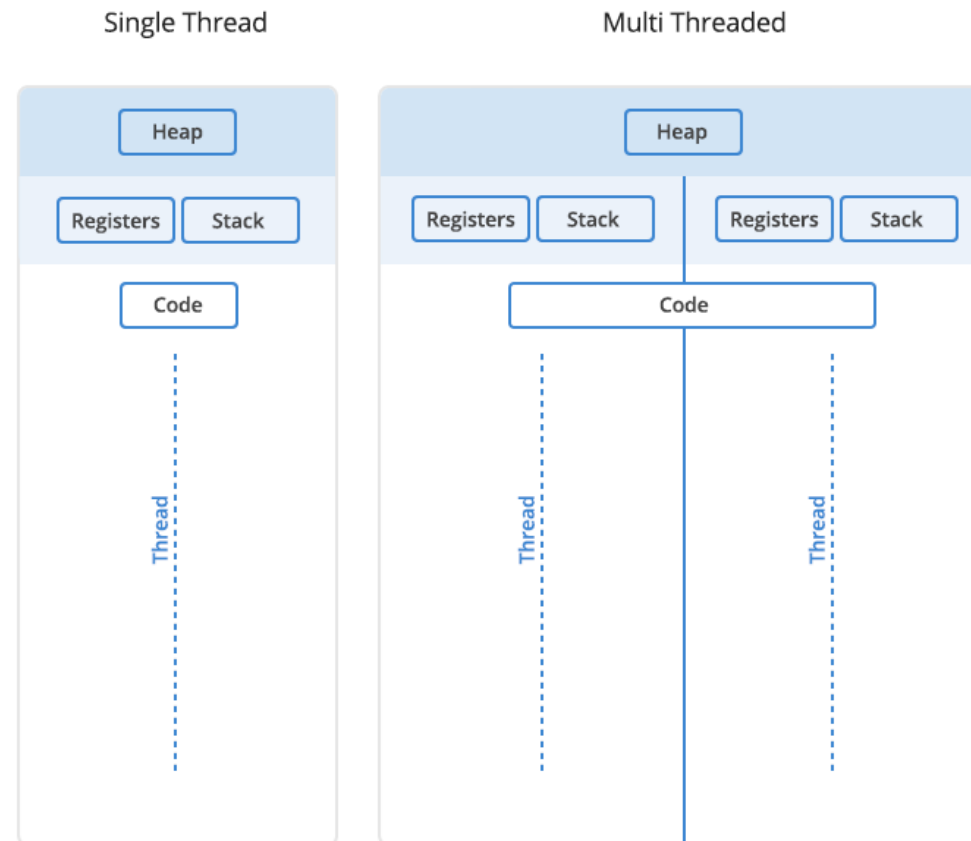
Processes vs Threads

Why use threads?

- Avoid needless blocking:
 - a single-threaded process will block when doing I/O; in a multi-threaded process, the operating system can switch the CPU to another thread in that process.
- Exploit parallelism:
 - the threads in a multi-threaded process can be scheduled to run in parallel on a multiprocessor or multicore processor.
- Avoid process switching:
 - structure large applications not as a collection of processes, but through multiple threads.

Processes vs Threads

- A thread could be defined as the unit of execution within a process.



Threads could be seen as lightweight processes as they have their own stack but can access shared data.

Which ones are more vulnerable to Problems ?



Processes vs Threads




Figure 3-3. A multithreaded server organized in a dispatcher/worker model.

Processes vs Threads

Design Choose: Process over Thread or Thread over Process?

Use-cases:

- The Chrome browser
- Spreadsheets
- A web client

Processes vs Threads

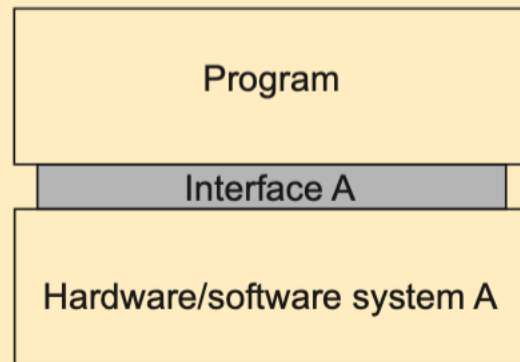
Design Choose: Process over Thread or Thread over Process?

Use-cases:

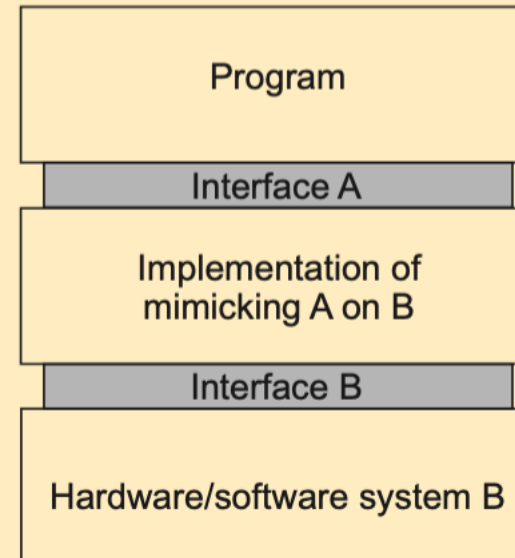
- The Chrome browser (multiple processes)
- Spreadsheets (multiple threads)
- A web client (multiple threads)

Virtualization

Principle: mimicking interfaces



(a)



(b)

Why do we need it?

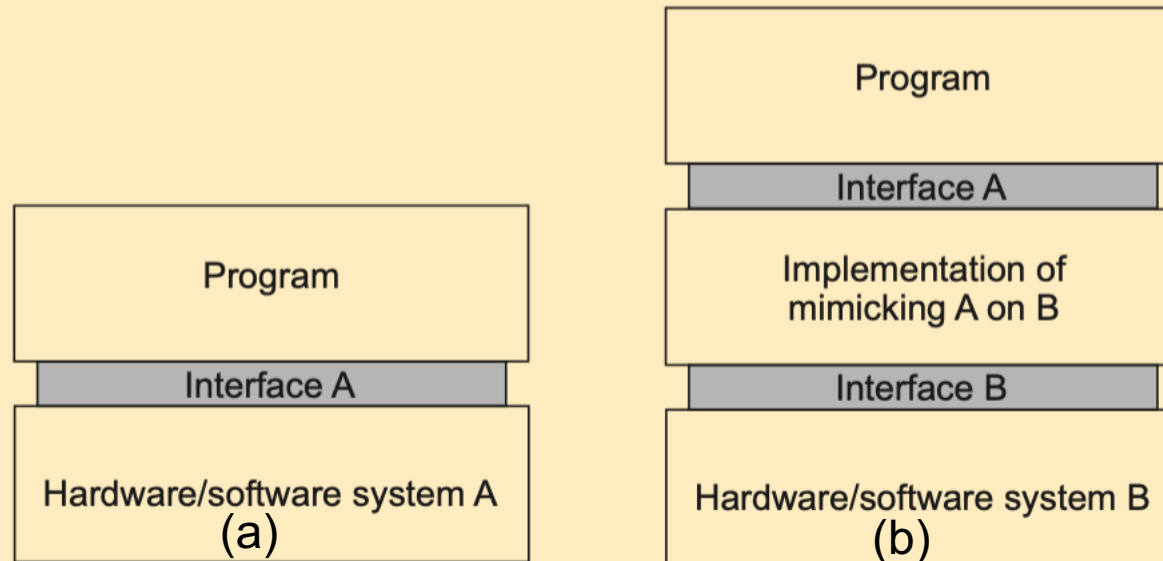


- (a) General organization between a program, interface, and system.
- (b) General organization of virtualizing system A on top of system B.

Virtualization

- Virtualization is important:
 - Hardware changes faster than software
 - Ease of portability and code migration,
 - Isolation of failing or attacked components

Principle: mimicking interfaces



(a) General organization between a program, interface, and system.

(b) General organization of virtualizing system A on top of system B.

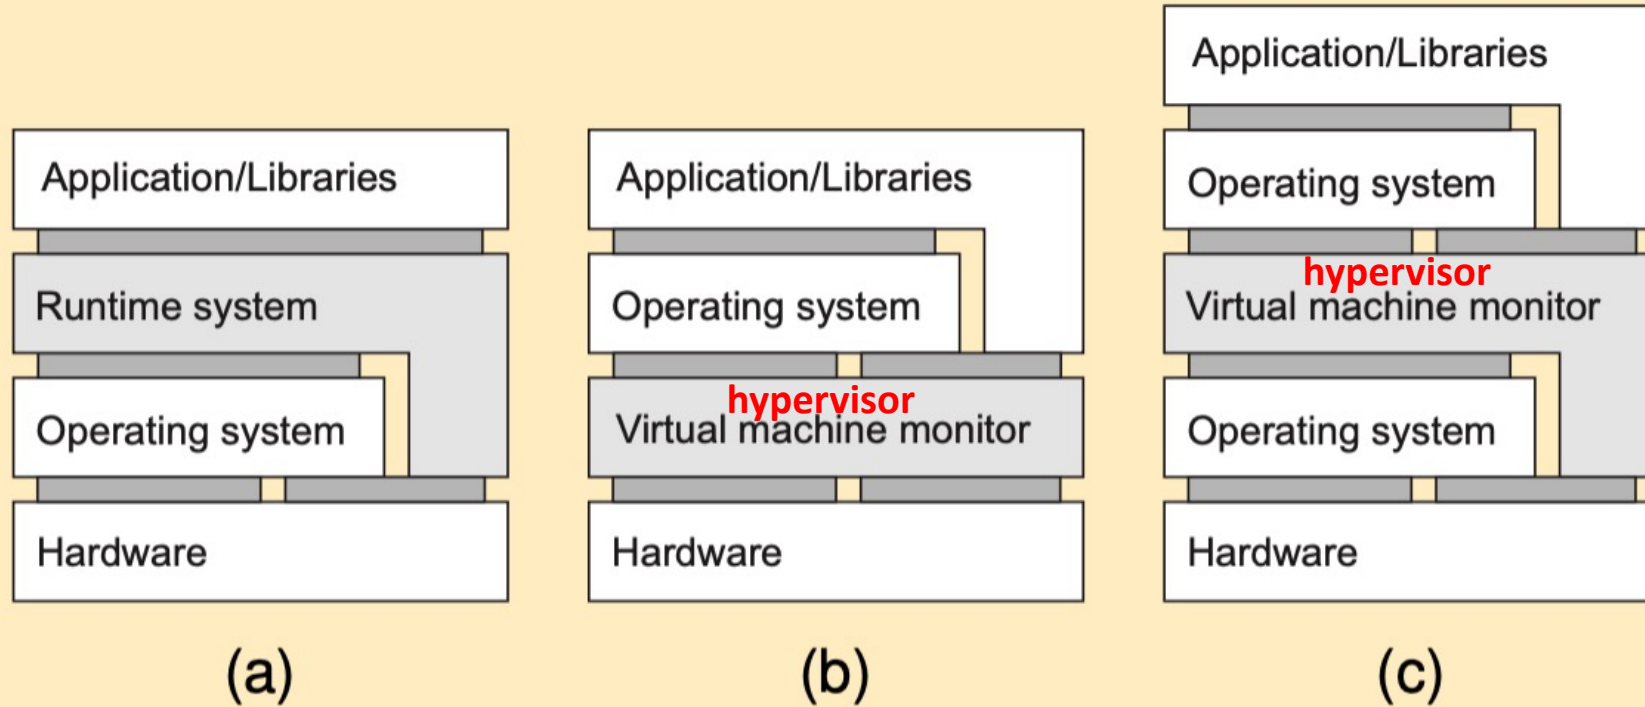
Virtualization

Mimicking interfaces

- Four types of interfaces at three different levels
 - **Instruction set architecture**: the set of *machine instructions*, with two subsets:
 - **Privileged instructions**: allowed to be executed only by the operating system.
 - **General instructions**: can be executed by any program.
 - **System calls** as offered by an operating system.
 - **Library calls**, known as an application programming interface (**API**)

Ways of virtualization


(a) Process VM, (b) Native VMM, (c) Hosted VMM



- (a) Separate set of instructions, an interpreter/emulator, running atop an OS.
- (b) Low-level instructions, along with bare-bones minimal operating system
- (c) Low-level instructions, but delegating most work to a full-fledged OS.

Zooming into VMs: performance

Refining the organization

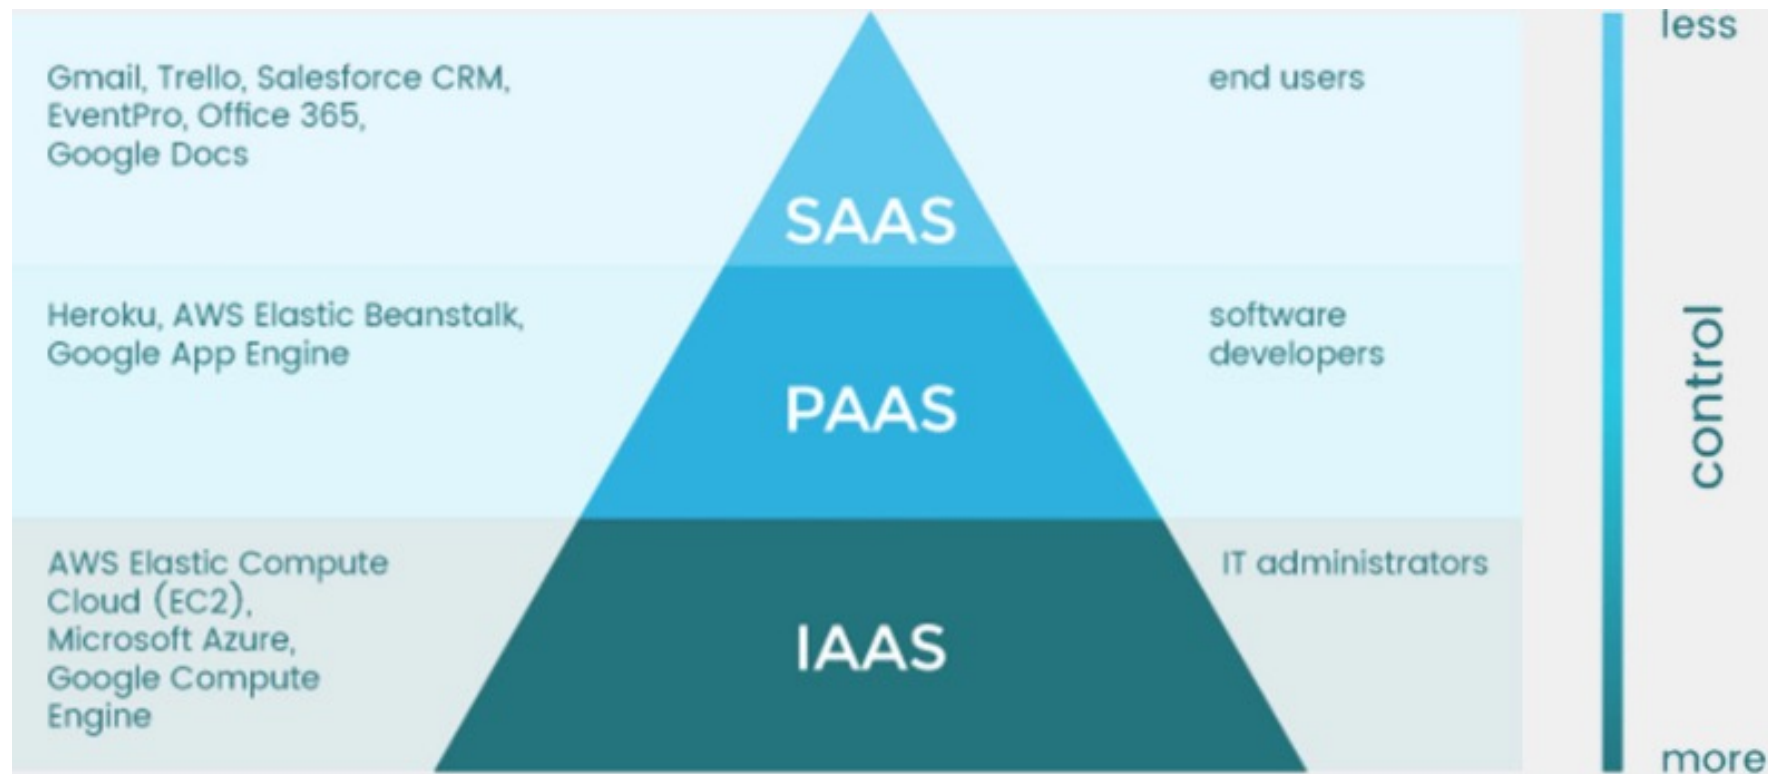


- Control-sensitive instruction: may affect configuration of a machine (e.g., one affecting relocation register or interrupt table).
- Behavior-sensitive instruction: effect is partially determined by context (e.g., POPF sets an interrupt-enabled flag, but only in system mode).

VMs and cloud computing

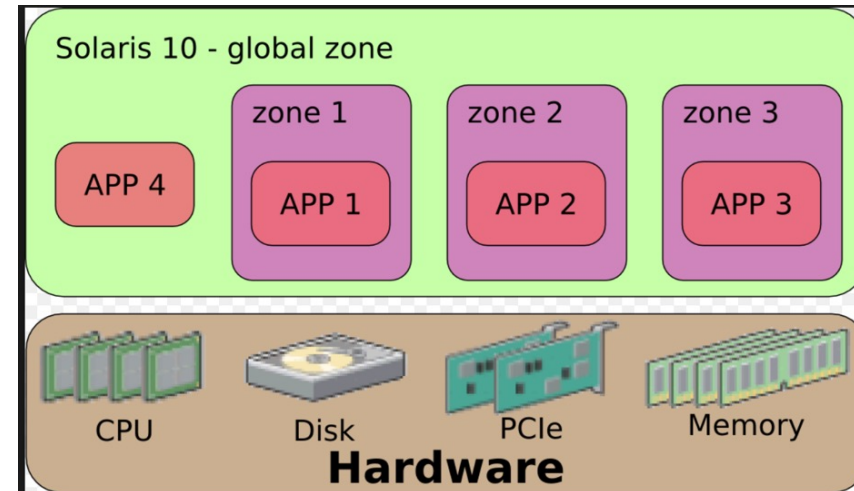
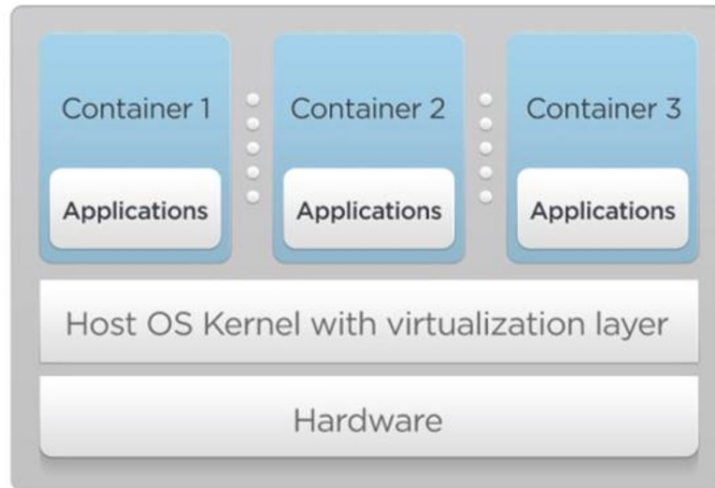
Three types of cloud services

- **Infrastructure-as-a-Service** covering the basic infrastructure
- **Platform-as-a-Service** covering system-level services
- **Software-as-a-Service** containing actual applications



containers

- Emulate OS-level interface with native interface
- “Lightweight” virtual machines
 - No **hypervisor**, OS provides necessary support

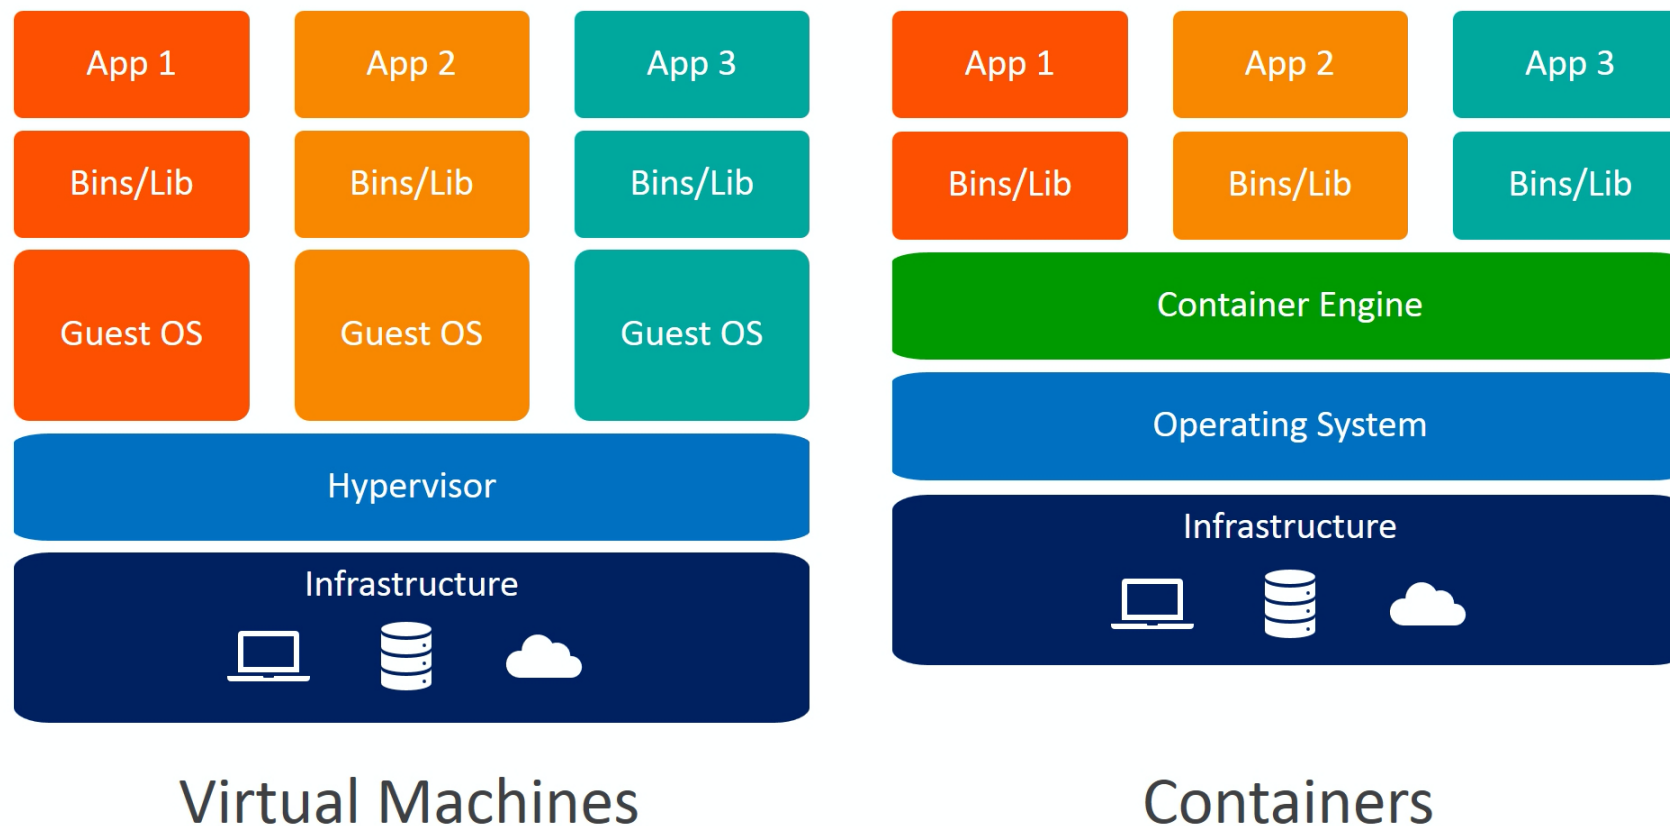


- Referred to as *containers*
 - Solaris containers, BSD jails, Linux containers

Docker and Linux Containers

- Linux containers are a set of kernel features
 - Need user space tools to manage containers
 - Virtuozzo, OpenVZm, VServer, Lxc-tools, Docker
- What does Docker add to Linux containers?
 - Portable container deployment across machines
 - Application-centric: geared for app deployment
 - Automatic builds: create containers from build files
 - Component re-use
- Docker containers are self-contained: no dependencies

VMs vs Containers



- Containers share OS kernel of the host
 - OS provides resource isolation
- Benefits
 - Fast provisioning, bare-metal like performance, lightweight

VMs vs Containers

VMs	Containers
Heavyweight	Lightweight
Limited performance	Native performance
Each VM runs in its own OS	All containers share the host OS
Hardware-level virtualization	OS virtualization
Startup time in minutes	Startup time in milliseconds
Allocates required memory	Requires less memory space
Fully isolated and hence more secure	Process-level isolation, possibly less secure

A To-Do List

- Read Chapters 3, Sec 1, 2 and part of 3
- Building your team