

COMP 474/6741 Intelligent Systems (Winter 2024)

Worksheet #6: Intelligent Agents

Task 1. What kind of question would you expect a *Concordia Chatbot* to be able to handle?

1. For *new students* (not currently registered at Concordia):

-
-
-

2. For *current students* (already enrolled at Concordia):

-
-
-

Task 2. Write a *regular expression* that matches different variations of naming <https://www.wikidata.org/wiki/Q326342>: “Concordia”, “Concordia U.”, “CU”, “Concordia University”, “Université Concordia”, ...:

Test it at <https://regex101.com/>

Task 3. Suppose we want to use an existing set of 1000 questions for training a ML classifier. If we use tf-idf vectors to represent each question, how many dimensions will the vectors have (make a rough estimate)?

Task 4. Ok, here is an (extremely simplified) idea of creating 2D feature vectors out of a natural language question: The first dimension *a* encodes the first occurrence of a question word (see table below) and the second dimension *b* the number of Capital Letters in the sentence:

Contains?	Value	#	Question	a	b	Class
Who	1	1	Where is Concordia?			Location
What	2	2	Who was Steve Jobs?			Definition
Where	3	3	What city is McGill in?			Location
(none)	0	4	What is NLP?			Definition



Task 5. Using the online parser at <https://corenlp.run/>, create a *parse tree* for the sentence *What is McGill?*. Note that you can now extract the *subject* of the sentence, e.g., to plug it into a SPARQL query.

Task 6. Now apply the kNN classification algorithm on the new question below to classify its type, according to the training data from Task 4. Use $k = 3$ and the Euclidian distance $d(\vec{p}, \vec{q}) = \sqrt{\sum_{i=1}^n (p_i - q_i)^2}$:

#	Question	a	b	d-Q1	d-Q2	d-Q3	d-Q4	Class?
5	What is McGill?							

You can now match the new question with a corresponding SPARQL template to obtain a query for your knowledge graph, filling in variables with the values extracted from the question.



Task 7. Now define a SPARQL template that can obtain information about a *person* from DBpedia. To keep it simple, for now we assume that the name extracted via NLP from the question is identical to the full name stored in the (English) `label` field for the subject (e.g., “Steve Jobs”).

```
PREFIX dbr: <http://dbpedia.org/resource/>
SELECT . . .
WHERE {

    . . .

}
```

Task 8. Create a *competency question* and a corresponding SPARQL query for our FOCU university example:

.....
Testing query for this question:

```
SELECT . . .
WHERE {

    . . .

}
```



Task 9. An early, well-known commercial service for semantic annotation of textual (mostly news) documents was Thompson Reuter’s *OpenCalais*, which has since been spun out and re-branded as *LSEG Data & Analytics PermID* (formerly Refinitif Intelligent Tagging): Try out the online demo at <https://permid.org/tagging> on a document, for example the first part of the Wikipedia article on Concordia. Look at the entities that were detected and go to the “RDF view”: what ID is given to Concordia in this knowledge graph?

.....
Hint: There is another tool at the top of the page, *Entity Search*, where you can cross-check your entities.



Task 10. Go to the DBpedia *Spotlight* online demo at <https://demo.dbpedia-spotlight.org>. Try analyzing a test document with some ambiguities, e.g., “*Paris Hilton went to the Hilton in Paris.*” Inspect the entities that were linked to DBpedia. Are they correct?

Task 11. For the questions in Task 1 above, which of the chatbot techniques covered so far would be able to answer them?